MATH 4322 Final Project Group 9

Introduction

Logistic Regression (Ryan Nguyen, Alan Johnson)

Paragraph explaining why we are using logistic regression models and the advantages and disadvantages of the model.

Model Formula

$$P(cardio = 1|X) = \frac{e^{(\beta_0+\beta_1 X_1+\beta_2 X_2+...+\beta_n X_n)}}{1 + e^{(\beta_0+\beta_1 X_1+\beta_2 X_2+...+\beta_n X_n)}}$$

Model 1 - Include all predictors

```
library(readr)
cardio_train <- read_delim("cardio_train.csv",
                           delim = ";", escape_double = FALSE, trim_ws = TRUE)
```

```
Rows: 70000 Columns: 13
-- Column specification --------------------------------------------------------
Delimiter: ";"
dbl (13): id, age, gender, height, weight, ap_hi, ap_lo, cholesterol, gluc, ...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
cardio_train$gender = as.factor(cardio_train$gender)
cardio_train$cholesterol = as.factor(cardio_train$cholesterol)
cardio_train$gluc = as.factor(cardio_train$gluc)
cardio_train$smoke = as.factor(cardio_train$smoke)
```

```r
cardio_train$alco = as.factor(cardio_train$alco)
cardio_train$active = as.factor(cardio_train$active)
cardio_train$cardio = as.factor(cardio_train$cardio)

heart.logistic1 = glm(cardio ~ . - id, family = "binomial",
                      data = cardio_train)
```

```
Warning: glm.fit: algorithm did not converge
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```r
summary(heart.logistic1)
```

```
Call:
glm(formula = cardio ~ . - id, family = "binomial", data = cardio_train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-8.4904  -0.9635  -0.0980   0.9907   4.6621

Coefficients:
               Estimate Std. Error z value Pr(>|z|)
(Intercept)  -8.084e+00  2.213e-01 -36.535  < 2e-16 ***
age           1.485e-04  3.557e-06  41.735  < 2e-16 ***
gender2       1.430e-02  2.107e-02   0.679    0.497
height       -5.626e-03  1.232e-03  -4.567 4.95e-06 ***
weight        1.521e-02  6.607e-04  23.023  < 2e-16 ***
ap_hi         3.951e-02  6.057e-04  65.235  < 2e-16 ***
ap_lo         3.004e-04  6.735e-05   4.460 8.18e-06 ***
cholesterol2  4.222e-01  2.593e-02  16.285  < 2e-16 ***
cholesterol3  1.134e+00  3.444e-02  32.929  < 2e-16 ***
gluc2         3.011e-02  3.438e-02   0.876    0.381
gluc3        -3.387e-01  3.809e-02  -8.894  < 2e-16 ***
smoke1       -1.314e-01  3.320e-02  -3.958 7.57e-05 ***
alco1        -1.695e-01  4.026e-02  -4.211 2.54e-05 ***
active1      -2.101e-01  2.105e-02  -9.981  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 97041  on 69999  degrees of freedom
Residual deviance: 80883  on 69986  degrees of freedom
AIC: 80911

Number of Fisher Scoring iterations: 25
```

Paragraph explaining which predictors are significant (look at significance table output)

Model 2 - Only include statistically significant predictors

```
heart.logistic2 =  glm(cardio  ~  age+height+weight+ap_hi+ap_lo+cholesterol+smoke+alco+act
                       ,  family  =  "binomial",
                     data  =  cardio_train)

summary(heart.logistic2)
```

```
Call:
glm(formula = cardio ~ age + height + weight + ap_hi + ap_lo +
    cholesterol + smoke + alco + active, family = "binomial",
    data = cardio_train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-8.4904  -0.9639  -0.0992   0.9900   4.6678

Coefficients:
               Estimate Std. Error z value Pr(>|z|)
(Intercept)  -8.124e+00  2.028e-01 -40.062  < 2e-16 ***
age           1.476e-04  3.551e-06  41.550  < 2e-16 ***
height       -5.356e-03  1.103e-03  -4.857 1.19e-06 ***
weight        1.516e-02  6.586e-04  23.023  < 2e-16 ***
ap_hi         3.960e-02  6.047e-04  65.485  < 2e-16 ***
ap_lo         3.028e-04  6.765e-05   4.475 7.63e-06 ***
cholesterol2  4.234e-01  2.497e-02  16.959  < 2e-16 ***
cholesterol3  9.855e-01  2.962e-02  33.275  < 2e-16 ***
smoke1       -1.222e-01  3.205e-02  -3.812 0.000138 ***
alco1        -1.641e-01  4.013e-02  -4.090 4.31e-05 ***
active1      -2.085e-01  2.104e-02  -9.909  < 2e-16 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 97041  on 69999  degrees of freedom
Residual deviance: 80984  on 69989  degrees of freedom
AIC: 81006

Number of Fisher Scoring iterations: 8
```

```
step(heart.logistic1)
```

```
Start:  AIC=80910.62
cardio ~ (id + age + gender + height + weight + ap_hi + ap_lo +
    cholesterol + gluc + smoke + alco + active) - id

              Df Deviance   AIC
- gender       1    80882 80908
<none>              80883 80911
- smoke        1    80900 80926
- alco         1    80902 80928
- ap_lo        1    80902 80928
- height       1    80917 80943
- gluc         2    80960 80984
- active       1    80987 81013
- weight       1    81248 81274
- cholesterol  2    82098 82122
- age          1    82467 82493
- ap_hi        1    87965 87991

Step:  AIC=80907.87
cardio ~ age + height + weight + ap_hi + ap_lo + cholesterol +
    gluc + smoke + alco + active

              Df Deviance   AIC
<none>              80882 80908
- ap_lo        1    80901 80925
- smoke        1    80901 80925
- alco         1    80901 80925
- height       1    80921 80945
- gluc         2    80984 81006
```

```
- active       1    80986 81010
- weight       1    81247 81271
- cholesterol  2    82099 82121
- age          1    82466 82490
- ap_hi        1    87982 88006


Call:  glm(formula = cardio ~ age + height + weight + ap_hi + ap_lo +
    cholesterol + gluc + smoke + alco + active, family = "binomial",
    data = cardio_train)

Coefficients:
 (Intercept)           age        height        weight         ap_hi
  -8.1440960     0.0001485    -0.0052551     0.0152121     0.0395279
       ap_lo  cholesterol2  cholesterol3         gluc2         gluc3
   0.0003007     0.4217513     1.1337397     0.0300489    -0.3389234
      smoke1         alco1       active1
  -0.1256165    -0.1681051    -0.2100651

Degrees of Freedom: 69999 Total (i.e. Null);  69987 Residual
Null Deviance:      97040
Residual Deviance: 80880    AIC: 80910
```

Model 3 - Using predictors from stepwise regression

Paragraph explaining the results of stepwise regression

```
heart.logistic3 = glm(formula = cardio ~ age + height + weight + ap_hi + ap_lo +
    cholesterol + gluc + smoke + alco + active, family = "binomial",
    data = cardio_train)
summary(heart.logistic3)
```

```
Call:
glm(formula = cardio ~ age + height + weight + ap_hi + ap_lo +
    cholesterol + gluc + smoke + alco + active, family = "binomial",
    data = cardio_train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-8.4904  -0.9638  -0.0979   0.9908   4.6623
```

```
Coefficients:
               Estimate Std. Error z value Pr(>|z|)
(Intercept)  -8.144e+00  2.030e-01 -40.123  < 2e-16 ***
age           1.485e-04  3.556e-06  41.760  < 2e-16 ***
height       -5.255e-03  1.104e-03  -4.760 1.94e-06 ***
weight        1.521e-02  6.607e-04  23.024  < 2e-16 ***
ap_hi         3.953e-02  6.051e-04  65.330  < 2e-16 ***
ap_lo         3.007e-04  6.736e-05   4.464 8.04e-06 ***
cholesterol2  4.218e-01  2.592e-02  16.273  < 2e-16 ***
cholesterol3  1.134e+00  3.444e-02  32.921  < 2e-16 ***
gluc2         3.005e-02  3.438e-02   0.874    0.382
gluc3        -3.389e-01  3.809e-02  -8.898  < 2e-16 ***
smoke1       -1.256e-01  3.209e-02  -3.914 9.07e-05 ***
alco1        -1.681e-01  4.021e-02  -4.181 2.90e-05 ***
active1      -2.101e-01  2.105e-02  -9.980  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 97041  on 69999  degrees of freedom
Residual deviance: 80882  on 69987  degrees of freedom
AIC: 80908

Number of Fisher Scoring iterations: 13
```

Determining Best Model

```r
extract_info  = function(model) {
  deviance <- summary(model)$null.deviance
  residual_deviance <- summary(model)$deviance
  r_squared <- 1 - (residual_deviance / deviance)
  AIC <- AIC(model)
  BIC <- BIC(model)

  return(c(Null_Deviance = deviance,
           Residual_Deviance = residual_deviance,
           R_Squared = r_squared,
           AIC = AIC,
           BIC = BIC))
}
```

```r
# Extract information from each model
info1 <- extract_info(heart.logistic1)
info2 <- extract_info(heart.logistic2)
info3 <- extract_info(heart.logistic3)

# Create a data frame to store the information
model_info <- data.frame(
  Model = c("heart.logistic1", "heart.logistic2", "heart.logistic3"),
  Null_Deviance = c(info1["Null_Deviance"], info2["Null_Deviance"], info3["Null_Deviance"]
  Residual_Deviance = c(info1["Residual_Deviance"], info2["Residual_Deviance"], info3["Res
  R_Squared = c(info1["R_Squared"], info2["R_Squared"], info3["R_Squared"]),
  AIC = c(info1["AIC"], info2["AIC"], info3["AIC"]),
  BIC = c(info1["BIC"], info2["BIC"], info3["BIC"])
)
(model_info)
```

```
            Model Null_Deviance Residual_Deviance R_Squared      AIC      BIC
1 heart.logistic1      97040.58          80882.62 0.1665072 80910.62 81038.81
2 heart.logistic2      97040.58          80983.71 0.1654656 81005.71 81106.43
3 heart.logistic3      97040.58          80881.87 0.1665149 80907.87 81026.91
```

Note: Model 2 removed gender and cholesterol and Model 3 just removed gender

Note: Model 3 has the lowest AIC and BIC

Final Equation for Logistic Regression Model

Insert latex equation here

Training/ Validation

```r
set.seed(100)
for(i in 1:10){
  # initialize vector to store prediction errors
test_errors = numeric(10)
sample= sample.int(n = nrow(cardio_train), size = floor(0.80*nrow(cardio_train)))
train.heart.logistic = cardio_train[sample,]
test.heart.logistic = cardio_train[-sample,]
```

```r
train.logistic = glm(formula = cardio ~ age + height + weight + ap_hi + ap_lo +
    cholesterol + gluc + smoke + alco + active, family = "binomial",
    data = cardio_train)

glm.pred = predict.glm(train.logistic, newdata = test.heart.logistic, type = "response")

# Convert probability to binary
test_predictions_binary = ifelse(glm.pred > 0.5, 1, 0)

# Calculate test prediction error
test_error= mean(test_predictions_binary != test.heart.logistic$cardio)

test_errors[i] = test_error
}

(mean_test_error = mean(test_errors))
```

```
[1] 0.02796429
```

Paragraph explaining the the procedure above and the mean error rate

Results

Insert graphics

Two paragraphs to provide the interpretation of results and your conclusions as it pertains to the original overall question.